

# High-Throughput Kafka-to-MySQL Processor with Redis Backup

## 1. Introduction

This document proposes the design of a **Delivery Tracking Microservice** to efficiently handle and persist a high volume of Kafka messages while minimizing the load on the MySQL database.

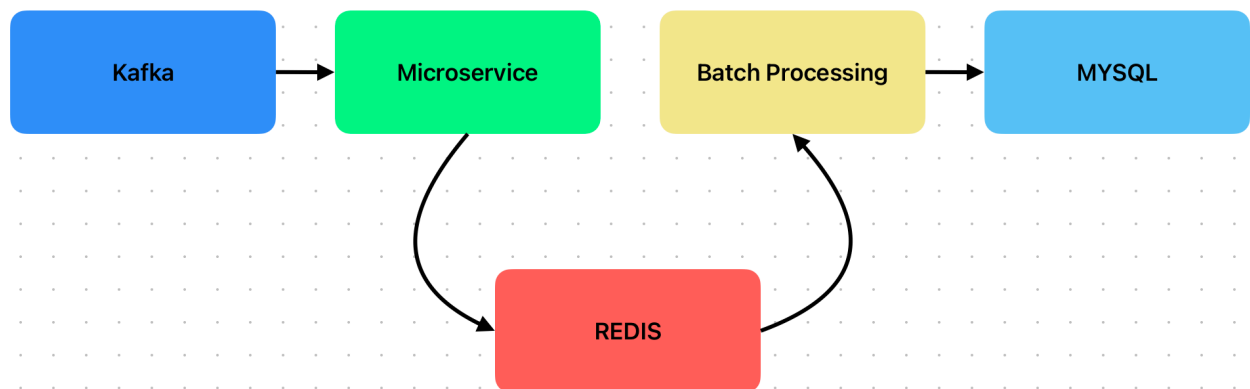
The core objective is to **reduce direct database hits** by introducing a buffering mechanism using **in-memory structures and Redis**, followed by **batch inserts** into MySQL at fixed intervals (approximately every 30 seconds).

### 1.1 Scope & Assumption

- Handle ~463 messages/sec across distributed instances (40 Million messages per day)
- Minimize DB hits using Redis-backed batching

## 2. System Architecture

### 2.1 High-Level Diagram



### 1. Kafka (Message Ingestion Layer)

- Kafka serves as the primary ingestion point, receiving messages related to delivery tracking in real time.
- The microservice subscribes to relevant Kafka topics to consume these messages.

### 2. Microservice (Processing Layer)

- A Spring Boot-based microservice that processes incoming Kafka messages.
- Messages are buffered **In-memory** and **Redis** for short-term storage.
- Ensures fault tolerance by persisting unprocessed messages in Redis in case of failures.

### 3. Redis (Temporary Buffer Storage)

- Acts as a distributed caching layer to hold messages temporarily.
- Ensures resilience and prevents data loss during unexpected failures.
- Supports quick retrieval for batch processing.

### 4. Batch Processing (Database Interaction Layer)

- Periodically (every 30 seconds), the batch processor fetches messages from **Redis and In-Memory buffers**.
- Performs batch inserts into MySQL, reducing individual database hits.
- Ensures efficient write operations, minimizing transaction overhead.

## 5. MySQL (Persistent Storage Layer)

- Stores the final processed delivery tracking data.
- Designed to support optimized bulk inserts for high-throughput scenarios.

# 3. Design Highlights

## 3.1 In-Memory Buffering

- Holds up to 1000 messages per pod
- Triggers batch insert either:
  - When buffer size hits 1000 or every 30 seconds

## 3.2 Redis Backup

- Key format: `msg:buffer:<instance-id>:<uuid>`

## 3.3 MySQL Batch Insert

- JDBC batch insert via `JdbcTemplate`
- Batches of up to 1000 rows

# 4. Crash Recovery

Failure	Solution
Pod crash (OOM etc.)	Restore from Redis on startup
MySQL outage	Retry with exponential backoff
Redis failure	Fallback to in-memory (data loss risk)
Message duplication	Use idempotent inserts

# 5. Redis & Memory Estimations

## 5.1 Message Size Calculation

- Each message: 500 characters

- Java String overhead: ~40 bytes per string
- Redis storage overhead: ~100 bytes per key-value pair
- Estimated size per message: ~700 bytes

## 5.2 Memory Requirements

### 5.2.1 For 40 million messages/day:

- Messages per second: ~463 messages/second (40M / 86400)
- Batch size: 1000 messages
- Batches per second: ~0.46 (every ~2 seconds)

### 5.2.2 In-Memory Buffer:

- Max buffer size: 1000 messages
- Memory needed:  $1000 * 700 \text{ bytes} = \sim 0.7 \text{ MB}$  per instance

### 5.2.3 Redis Requirements:

- Worst case (all messages in Redis):  $40\text{M} * 700 \text{ bytes} = \sim 28\text{GB}$
- Practical case (buffer overflow only):
  - If we have 10 consumer instances
  - Each might have up to 1000 messages in Redis
  - Total:  $10 * 1000 * 700 \text{ bytes} = \sim 7\text{MB}$

## 5.3 Recommended Redis Sizing

- Start with 1GB Redis instance
- Enable Redis persistence for crash recovery
- Consider Redis Cluster if high availability is critical

## 6. Conclusion

- ✓ High Throughput (40M messages/day)
- ✓ Fault Tolerance with Redis
- ✓ Optimized MySQL inserts
- ✓ Kubernetes Ready
- ✓ Resilient and Observable